

COSC2406/2407 Database Systems

Tree Index Structures

Assignment Project Exam Help

Xiangmin (Emily) Zhou

Only

<https://eduassistpro.github.io>

Thursdays

Email : xiangmin.zhou@rm

Add WeChat [edu_assist_pr](#)

Lecture 5

References: Ramakrishnan & Gehrke Chapter 10

Garcia-Molina et al. Chapter 13

Elmasri & Navathe Chapter 5

Diagrams courtesy Ramakrishnan & Gehrke

Assignment Project Exam Help

In this lect

Specific

- The <https://eduassistpro.github.io>
- Dynamic B+-trees

Add WeChat edu_assist_pr

- As for any index, 3 alternatives for index data entries k^* :

- 1 Data record with search key value k
- 2 $(k, \text{rid of data record with search key value } k)$
- 3

– Ch


- Tre and equality searches
- ISAM (Indexed Sequential Access Meth structure
- B+-tree is a dynamic structure that adjusts w deletions

Motivation for Tree Structures: Range Searches

- *"Find all students older than 22."*

If data is in sorted file, we can binary search to find first such student, then scan to find other matches


- The cost of binary search can be high
- Sim



- Now we can binary search on the (smaller) index file!

Motivation for Tree Structures: Range Searches ...

Index file may still be quite large. Apply idea repeatedly.



Leaf pages contain data entries.

Assignment Project Exam Help

Consider a heap file, with 2 records per page:

<https://eduassistpro.github.io>

Data
Pages

2*	5*	27*	55*	3*	37*	51*	97*
----	----	-----	-----	----	-----	-----	-----

Add WeChat edu_assist_pro

Assignment Project Exam Help

Sort file on search key:

<https://eduassistpro.github.io>

Data
Pages

1*	5*	20*	27*	33*	37*	40*		97*
----	----	-----	-----	-----	-----	-----	--	-----

Add WeChat edu_assist_pro

Build an index on file: **Assignment Project Exam Help**

<https://eduassistpro.github.io>



Build a sparse index on file:

Assignment Project Exam Help

<https://eduassistpro.github.io>



Build an index on the index:


Assignment Project Exam Help

<https://eduassistpro.github.io>

Index

Add WeChat edu_assist_pr

Data
Pages



The tree we have constructed is a so-called ISAM (Indexed Sequential Access Method) structure. (See Section 10.2)

Each index page has *index data entries* of the form (key, pointer).


Each ind

A key serv

pointed t

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



- File creation: Leaf (data) pages allocated sequentially, sorted by search key; then index pages allocated, then space for overflow pages.

- Ind

sea

- Se

Co

\log_F

N is


number of leaf pages.

- Insert: Find the leaf page where the data entry belongs and insert it there. Create overflow pages if necessary.


- Delete: Find and remove the entry from the leaf; if this empties an overflow page, de-allocate the page.

Example ISAM Tree

Each node can hold 2 entries no need for 'next-leaf-page' pointers
(because of the sequential allocation of leaf pages).




After Inserting 23*, 48*, 41*, 42* ...



Advantage: Less locking problems in ISAM than in other structures, since no index page is changed.

... Then Deleting 42*, 51*, 97*



Note that 51* appears in index levels, but not in a leaf!

Assignment Project Exam Help

Once an ISAM file is created, inserts and deletes affect only the leaf page content.

As a result

required to
well.

To reduce

(20%, say) for future insertions.

Otherwise, the only way to get rid of overflow chain
reorganisation of the whole file structure.

<https://eduassistpro.github.io>

Add WeChat [edu_assist_pro](#)

Assignment Project Exam Help

The fact that only leaf pages are adjusted gives a big advantage for *concurrent* access:

When a page is updated, only the leaf page is modified, not the index pages.

concurrent

of users will

significa

the root of a tree.

Since ISAM index pages are never modified, they are never locked.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

B+-tree: The Most Widely Used Index

The B+-tree is the most widely used index structure (see Section 10.3). It has the following characteristics:


- Internal nodes direct the search, index data entries are in the leaf
- The
- Ins
nu
- Min
- Each node contains $d \leq m \leq 2$
 d is called the order of the tree.
- Supports equality and range-searches ef
- Leaf pages are organised as a double linked list for fast traversal and reorganisation

Assignment Project Exam Help

Index

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Non-leaf nodes contain m index entries, with $m + 1$ pointers to children (like SAM structure).

Leaf nodes contain data entries, either:

- act
ind
dat
- pointers to data records elsewhere on disk (the B+-tree is an *index file*, distinct from records)

If the indexed field is of fixed-length, so is the index entry. If the indexed field is variable-length, so are the index entries.


Example B+-tree

Assignment Project Exam Help

- A search begins at the root, and key comparisons direct the search to a leaf (as in ISAM)
- (Tr

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



- Find correct leaf L
- Put data entry into L
 - If L has enough space, done!
 -



<https://eduassistpro.github.io>

- This can happen recursively.
 - To split index node, redistribute entries e but *push-up* middle key. (Contrast role)
- Splits “grow” tree; root split increases height
 - Tree growth: gets wider or one level taller at top.

Inserting 8* into Example B+-tree

- Observe how minimum occupancy level is guaranteed in both leaf and index pages splits.

- No difference between *copy-up* and *push-up*!



Assignment Project Exam Help



- Notice that root was split, leading to increase in the number of leaf nodes.
- In this example, we can avoid split by re-distributing the keys, however, this is usually not done in practice.

Deleting a Data Entry from a B+-Tree

- Start at root; find the leaf L where the entry belongs.
- Remove the entry.
 - If L is at least half full, done!
-


<https://eduassistpro.github.io>

- If merge occurred, must delete entry (point x) from parent of L .
- Merge could propagate to root, decreasing height.

Refer to Section 10.6 for an example.

Garcia-Molina et al. illustrates a B-tree example in Section 13.3.6.

Assignment Project Exam Help



Deletion of 20* is solved by re-distribution from sib
the new splitting (middle) key is copied up.

Deleting 24*


Leaf nodes must be merged:

Assignment Project Exam Help

<https://eduassistpro.github.io>

Then, index nodes need to be merged and entry 1

Add WeChat edu_assist_pr



Assignment Project Exam Help

- Typical order (minimum node size): 100. Typical fill-factor: 67%.
 - Average fan-out = 133
- Ty
 - <https://eduassistpro.github.io>
- Can often hold top levels in the buffer manag
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 Mbytes

Assignment Project Exam Help

Read more indexes in MongoDB (including use of B+-trees)

<https://>

<https://eduassistpro.github.io/>

Read more about how B+-trees are implemented in Derby:

<http://db.apache.org/derby/papers/>

Add WeChat edu_assist_pr

Assignment Project Exam Help

We have
Specific
(ISAM) a

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr